

Parallel processor architecture & algorithms

Intel 8086 (1981) : 5 MHz used in first IBM PC

80486 (1989) : 25 MHz

→ i486 because of a court ruling that prohibited the trademarking of numbers

Pentium (1993) : 66 MHz

Pentium 4 (2000) : 1.5 GHz deep \approx 30-stage pipeline


Pentium D (2005) : 3.2 GHz

and then clock speed stopped increasing!

Quad core Xeon (2008) : 3 GHz


Key to performance scaling : increase the number of cores on chip.

Processors need data to compute on.




Problem: SRAM cannot support more than ~4 memory requests in parallel.

\$: cache P: processor




Most of the time program running on the processor accesses local memory or "cache" memory.



Every once in awhile, it accesses remote memory

Round-trip required

Research Idea: Execution Migration



When program running on a processor needs to access cache memory of another processor, it migrates its "context" to the remote processor & executes there.

One-way trip for data access

Context = Program Counter + Register File + ..
(can be larger than data to be accessed) few kbits

Assume we know or can predict the access pattern of a program

m_1, m_2, \dots, m_N memory addresses
 $p(m_1), p(m_2), \dots, p(m_N)$ processor caches for each m_i

Example: $P_1 P_2 P_2 P_1 P_1 P_3 P_2$

$$\text{cost}_{\text{mig}}(s, d) = \text{distance}(s, d) + L \leftarrow \begin{array}{l} \text{load latency} \\ \text{is a function} \\ \text{of context size.} \end{array}$$


$$\text{cost}_{\text{access}}(s, d) = 2 * \text{distance}(s, d)$$

if $s == d$, costs are defined to be 0.

Problem: Decide when to migrate to minimize total memory cost of trace.

Example:

Start at P_1



What can we use to solve this problem?

Dynamic Programming!

Program at P_i , initially, number of processors = Q

Subproblems?

$DP(k, p_i)$ = cost of optimal solution for the prefix $m_1 \dots m_k$ of memory accesses when program starts at P_i and ends up at p_i

$$DP(k+1, p_j) = \begin{cases} DP(k, p_j) + \text{cost access}(p_j, p(m_{k+1})) \\ \quad \text{if } p_j \neq p(m_{k+1}) \\ \min_{i=1}^Q (DP(k, p_i) + \text{cost mig}(p_i, p_j)) \\ \quad \text{if } p_j = p(m_{k+1}) \end{cases}$$

Complexity?

$$O(N \cdot Q)$$

number of subproblems

$\xrightarrow{\text{cost per subproblem}}$

$$= O(NQ^2)$$

My research group is building a 128-processor Execution Migration Machine that uses a migration predictor based on this analysis.


Erik's main research areas:

- computational geometry [6.850]
- geometric folding algorithms [6.849]
- self-assembly
- data structures [6.851]
- graph algorithms [6.889]
- recreational algorithms [SP.268]
- algorithmic sculpture

Geometric folding algorithms: [6.849, videos online]

- design: algorithms to fold any polyhedral surface from a square of paper
[Demaine, Demaine, Mitchell 2000; Demaine & Tachi 2011]
 - bicolor paper \Rightarrow can 2-color faces
 - OPEN: how to best optimize "scale factor"
 - e.g. best $n \times n$ checkerboard folding
recently improved from $\sim n/2$ \rightarrow $\sim n/4$
- foldability: given a crease pattern, can you fold it flat?
 - NP-complete in general [Bern & Hayes 1996]
 - OPEN: $m \times n$ map with creases specified as mountain/valley [Edmonds 1997]
 - just solved: $2 \times n$ [Demaine, Liu, Morgan 2011]

- hyperbolic paraboloid [Bauhaus 1929]
doesn't exist!




[Demaine, Demaine, Hart, Price, Tachi 2009]

- understanding circular creases

- any straight-line graph can be made by folding flat & one straight cut
[Demaine, Demaine, Lubiw 1998;
Bern, Demaine, Eppstein, Hayes 1999]

Self-assembly: geometric model of computation

- glue: e.g. DNA strands, each pair has strength
- square tiles with glue on each side
- Brownian motion: tiles/constructions stick together if \sum glue strengths \geq temperature



- can build $n \times n$ square using $O(\frac{\lg n}{\lg \lg n})$ tiles
[Rothenmund & Winfree 2000]

or using $O(1)$ tiles & $O(\lg n)$ "stages"
algorithmic steps by the bioengineer

[Demaine, Demaine, Fekete, Ishaque, Rafalin, Schweller, Souvaine 2007]

- can replicate ∞ copies of given unknown shape using $O(1)$ tiles & $O(1)$ stages


[Abel, Benbernou, Damian, Demaine, Demaine, Flatland, Komihers, Schweller 2010]

Data structures: [6.851, videos next semester]

- integer data structures: store n integers in $\{0, 1, \dots, u-1\}$ subject to insert, delete, predecessor, successor (on word RAM)
 - hashing does exact search in $O(1)$
 - AVL trees do all in $O(\lg n)$
 - $O(\lg \lg u)$ /op. [van Emde Boas]
 - $O(\frac{\lg n}{\lg \lg u})$ /op. [fusion trees:
Fredman & Willard]
 - $O(\sqrt{\frac{\lg n}{\lg \lg n}})$ /op. [min of above]

- cache-efficient data structures:

- memory transfers happen in blocks




- searching takes $\Theta(\log_B N)$ transfers (vs. $\lg n$)
- sorting takes $\Theta(\frac{N}{B} \log_C \frac{N}{B})$ transfers
- possible even if you don't know B & C !

(Almost) planar graphs: [6.889, videos online]

- Dijkstra in $O(n)$ time
[Henzinger, Klein, Rao, Subramanian 1997]
- Bellman-Ford in $O(n \lg^2 n / \lg \lg n)$ time
[Mozes & Wolff-Nilson 2010]

- many problems NP-hard, even in planar graphs
- but can find a solution within $1 + \varepsilon$ factor of optimal, for any ε
- run BFS from any root vertex r
- delete every k layers
- for many problems, solution messed up by only $1 + \frac{1}{k}$ factor ($\Rightarrow k = \frac{1}{\varepsilon}$)
- connected components of remaining graph have $< k$ layers ~ can solve via DP typically in $\sim 2^k \cdot n$ time



[Baker 1994 & others]

Recreational algorithms:

- many algorithms & complexities of games
[some in SP. 268 & our book
Games, Puzzles, & Computation (2009)]
- $n \times n \times n$ Rubik's Cube diameter is $\mathcal{O}(n^2/\lg n)$
[Demaine, Demaine, Eisenstat, Lubiw, Winslow 2011]
- Tetris is NP-complete
[Brenkelaar, Demaine, Hohenberger, Hoogeboom, Kosters, Liben-Nowell 2004]
- balloon twisting any polyhedron
[Demaine, Demaine, Hart 2008]
- algorithmic magic tricks

Algorithms classes at MIT: (post-6.006)

- #1: 6.046: Intermediate Algorithms
(more adv. algorithms & analysis, less coding)
- 6.047: Computational Biology
(genomes, phylogeny, etc.)
- 6.854: Advanced Algorithms
(intense survey of whole field)
- 6.850: Geometric Computing
(working with points, lines, polygons, meshes, ...)
- 6.849: Geometric Folding Algorithms
(origami, robot arms, protein folding, ...)
- 6.851: Advanced Data Structures
(sublogarithmic performance)
- 6.852: Distributed Algorithms
(reaching consensus in a network with faults)
- 6.853: Algorithmic Game Theory
(Nash equilibria, auction mechanism design, ...)
- 6.855: Network Optimization
(optimization in graph: beyond shortest paths)
- 6.856: Randomized Algorithms
(how randomness makes algs. simpler & faster)
- 6.857: Network and Computer Security
(cryptography)

Other theory classes:

- 6.045: Automata, Computability, & Complexity
- 6.840: Theory of Computing
- 6.841: Advanced Complexity Theory
- 6.842: Randomness & Computation
- 6.845: Quantum Complexity Theory
- 6.440: Essential Coding Theory
- 6.441: Information Theory

Top 10 Uses of 6.006 Cushions

10. Sit on it: guaranteed inspiration in constant time
(bring it to the final exam)
9. Frisbee (after cutting it into a circle)*
8. Sell as a limited-edition collectible on eBay
(they'll probably never be made again—at least \$5)
7. Put two back-to-back to remove branding*
(so no one will ever know you took this class)
6. Holiday conversation starter... and stopper
(we don't recommend re-gifting)
5. Asymptotically optimal acoustic paneling
(for practicing piano & guitar fingering DP)
4. Target practice for your next LARP*
(Live Action Role Playing)
3. Ten years from now, it might be all you'll remember about 6.006
(maybe also this top ten list)
2. Final exam cheat sheet*
1. *Three words:* OkCupid profile picture